

```
# importing necessary libraries

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# loading the dataset

crop_data=pd.read_csv("/content/crop_production.csv")
crop_data
```

↗

	State_Name	District_Name	Crop_Year	Season	Crop	Area	Production
0	Andaman and Nicobar Islands	NICOBARS	2000	Kharif	Arecanut	1254.0	2000.0
1	Andaman and Nicobar Islands	NICOBARS	2000	Kharif	Other Kharif pulses	2.0	1.0
2	Andaman and Nicobar Islands	NICOBARS	2000	Kharif	Rice	102.0	321.0
3	Andaman and Nicobar Islands	NICOBARS	2000	Whole Year	Banana	176.0	641.0
4	Andaman and Nicobar Islands	NICOBARS	2000	Whole Year	Cashewnut	720.0	165.0
...	...	...	...	...	...	...	...
103163	Madhya Pradesh	BALAGHAT	2000	Rabi	Safflower	6.0	1.0
103164	Madhya Pradesh	BALAGHAT	2000	Rabi	Wheat	14004.0	9796.0
103165	Madhya Pradesh	BALAGHAT	2000	Whole Year	Coriander	291.0	65.0
103166	Madhya Pradesh	BALAGHAT	2000	Whole Year	Dry chillies	405.0	72.0
103167	Madhya Pradesh	BALAGHAT	2000	Whole Year	Garlic	131.0	449.0

103168 rows × 7 columns

```
crop_data.shape

#rows X columns
```

↗ (103168, 7)

```
# dataset columns
crop_data.columns

↗ Index(['State_Name', 'District_Name', 'Crop_Year', 'Season', 'Crop', 'Area',
        'Production'],
        dtype='object')
```

```
# statistical inference of the dataset

crop_data.describe()
```


↗

	Crop_Year	Area	Production
count	103168.000000	103168.000000	1.009830e+05
mean	2005.893455	9081.339826	1.196607e+06
std	4.931049	30605.983819	2.528363e+07
min	1997.000000	0.040000	0.000000e+00
25%	2002.000000	87.000000	1.000000e+02
50%	2006.000000	566.000000	7.760000e+02
75%	2010.000000	3652.000000	6.771500e+03
max	2014.000000	877029.000000	1.125000e+09


```
# Checking missing values of the dataset in each column
crop_data.isnull().sum()
```

↗ State\_Name 0  
District\_Name 0  
Crop\_Year 0  
Season 0  
Crop 0  
Area 0  
Production 2185  
dtype: int64


```
# Dropping missing values
crop_data = crop_data.dropna()
crop_data
```



	State_Name	District_Name	Crop_Year	Season	Crop	Area	Production
0	Andaman and Nicobar Islands	NICOBARS	2000	Kharif	Arecanut	1254.0	2000.0
1	Andaman and Nicobar Islands	NICOBARS	2000	Kharif	Other Kharif pulses	2.0	1.0
2	Andaman and Nicobar Islands	NICOBARS	2000	Kharif	Rice	102.0	321.0
3	Andaman and Nicobar Islands	NICOBARS	2000	Whole Year	Banana	176.0	641.0
4	Andaman and Nicobar Islands	NICOBARS	2000	Whole Year	Cashewnut	720.0	165.0
...	...	...	...	...	...	...	...




```
#checking
crop_data.isnull().values.any()
```



False


```
# Displaying State Names present in the dataset
crop_data.State_Name.unique()
```




array(['Andaman and Nicobar Islands', 'Andhra Pradesh', 'Arunachal Pradesh', 'Assam', 'Bihar', 'Chandigarh', 'Chhattisgarh', 'Dadra and Nagar Haveli', 'Goa', 'Gujarat', 'Haryana', 'Himachal Pradesh', 'Jammu and Kashmir ', 'Jharkhand', 'Karnataka', 'Kerala', 'Madhya Pradesh'], dtype=object)

```
# Adding a new column Yield which indicates Production per unit Area.
```

```
crop_data['Yield'] = (crop_data['Production'] / crop_data['Area'])
crop_data.head(10)
```



	State_Name	District_Name	Crop_Year	Season	Crop	Area	Production	
0	Andaman and Nicobar Islands	NICOBARS	2000	Kharif	Arecanut	1254.0	2000.0	1
1	Andaman and Nicobar Islands	NICOBARS	2000	Kharif	Other Kharif pulses	2.0	1.0	0
2	Andaman and Nicobar Islands	NICOBARS	2000	Kharif	Rice	102.0	321.0	3
3	Andaman and Nicobar Islands	NICOBARS	2000	Whole Year	Banana	176.0	641.0	3
4	Andaman and Nicobar	NICOBARS	2000	Whole ..	Cashewnut	720.0	165.0	0



```
# Dropping unnecessary columns
```


```
data = crop_data.drop(['State_Name'], axis = 1)
```

```
import seaborn as sns
import matplotlib.pyplot as plt

# Select only numeric columns from your DataFrame
numeric_data = data.select_dtypes(include='number')

# Compute the correlation matrix for numeric data
corr_matrix = numeric_data.corr()

# Plotting the heatmap with annotations and a title
plt.figure(figsize=(10, 8)) # Adjust the figure size if needed
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix')
plt.show()
```



```
dummy = pd.get_dummies(data)
dummy
```

	Crop_Year	Area	Production	Yield	District_Name_AGAR MALWA	District_Name_Al
0	2000	1254.0	2000.0	1.594896		False
1	2000	2.0	1.0	0.500000		False
2	2000	102.0	321.0	3.147059		False
3	2000	176.0	641.0	3.642045		False
4	2000	720.0	165.0	0.229167		False
...	...	...	...	...		...
103163	2000	6.0	1.0	0.166667		False
103164	2000	14004.0	9796.0	0.699514		False
103165	2000	291.0	65.0	0.223368		False
103166	2000	405.0	72.0	0.177778		False
103167	2000	131.0	449.0	3.427481		False

100983 rows x 392 columns

```

from sklearn.model_selection import train_test_split


x = dummy.drop(["Production", "Yield"], axis=1)
y = dummy["Production"]

# Splitting data set - 25% test dataset and 75%

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=5)

print("x_train :", x_train.shape)
print("x_test :", x_test.shape)
print("y_train :", y_train.shape)
print("y_test :", y_test.shape)


```

 x\_train : (75737, 390)  
x\_test : (25246, 390)  
y\_train : (75737,)  
y\_test : (25246,)

```

print(x_train)
print(y_train)

```



59924	False	False	False	False
5560	False	False	False	False
20536	False	False	False	False
18709	False	False	False	False
35767	False	False	False	False
...	...	...	...	...
98859	False	False	False	False
40637	False	False	False	False
85066	False	False	False	False
87863	False	False	False	False
82926	False	False	False	False
...	...	...	...	...
59924	False	False	False	False
5560	False	False	False	False
20536	False	False	False	False
18709	False	False	False	False
35767	False	False	False	False
...	...	...	...	...
98859	False	False	False	False
40637	False	False	False	False
85066	False	False	False	False
87863	False	False	False	False
82926	False	False	False	False
...	...	...	...	...
59924	False	False	False	False
5560	False	False	False	False
20536	False	False	False	False
18709	False	False	False	False
35767	False	False	False	False
...	...	...	...	...
98859	False	False	False	False
40637	False	False	False	False
85066	False	False	False	False
87863	False	False	False	False
82926	False	False	False	False
...	...	...	...	...
59924	False	False	False	False
5560	False	False	False	False
20536	False	False	False	False
18709	False	False	False	False
35767	False	False	False	False
...	...	...	...	...
98859	False	False	False	False
40637	False	False	False	False
85066	False	False	False	False
87863	False	False	False	False
82926	False	False	False	False
...	...	...	...	...
59924	False	False	False	False
5560	False	False	False	False
20536	False	False	False	False
18709	False	False	False	False
35767	False	False	False	False

[75737 rows x 390 columns]  
98859 108.0  
40637 6.0  
85066 12745.0  
87863 2284.0  
82926 161.0  
...  
59924 1500.0  
5560 600.0  
20536 27.0  
18709 4779.0  
35767 81.0  
Name: Production, Length: 75737, dtype: float64

```
# Training the Simple Linear Regression model .
```

```

from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(x_train, y_train)

```

LinearRegression  
LinearRegression()

```
# Predicting the test Results
```

```
lr_predict = model.predict(x_test)
lr_predict
```

```
array([ -179651.33984375,  -989905.37890625, -3914674.02148438, ...,
        -5991530.2265625 ,  1107923.00585938, -1027953.0390625  ])
```

```
model.score(x_test,y_test)
```


```
0.21903980552344915
```

```
from sklearn.metrics import r2_score
r = r2_score(y_test,lr_predict)
print("R2 score : ",r)
```

```
R2 score : 0.21903980552344915
```

```
plt.scatter(y_test,lr_predict)
plt.xlabel('Actual')
plt.ylabel('Predicted')
plt.title('Linear Regression')
```

```
Text(0.5, 1.0, 'Linear Regression')
```



**Clearly, the dataset is not good for linear regression.**

## ✓ Using the random Forest regressor

```
from sklearn.ensemble import RandomForestRegressor
model = RandomForestRegressor(n_estimators = 11)
model.fit(x_train,y_train)
rf_predict = model.predict(x_test)
rf_predict
```

```
array([ 277.18181818, 2317.18181818,  15.54545455, ..., 1469.72727273,
        679.72727273, 2196.          ])
```

```
model.score(x_test,y_test)
```

```
0.9811208481033443
```

## ✓ Using Decision tree

```
# Training model
from sklearn.tree import DecisionTreeRegressor
regressor = DecisionTreeRegressor(random_state = 5)
regressor.fit(x_train,y_train)
```

```
# Predicting results
decisiontree_predict = regressor.predict(x_test)
decisiontree_predict
```

```
array([ 250., 3000.,    9., ..., 1442.,  473., 2080.])
```

```
regressor.score(x_test,y_test)
```

```
0.9753368569952419
```